

# HW2 Tutorial

# Logistics

# Logistics

Deliverables and Due Date:



**HW2 Part 1:** Implementing CNN from Scratch [7 points]

**HW2 Part 2:** PyTorch - Implement Two-Layer Network, Vanilla CNN  
and Your own model [5+2 points]

**Report:** [6 points]



**Due Date:** 11:59 PM Feb 19, 2024

**Grace Period:** 11:59 PM Feb 21, 2024

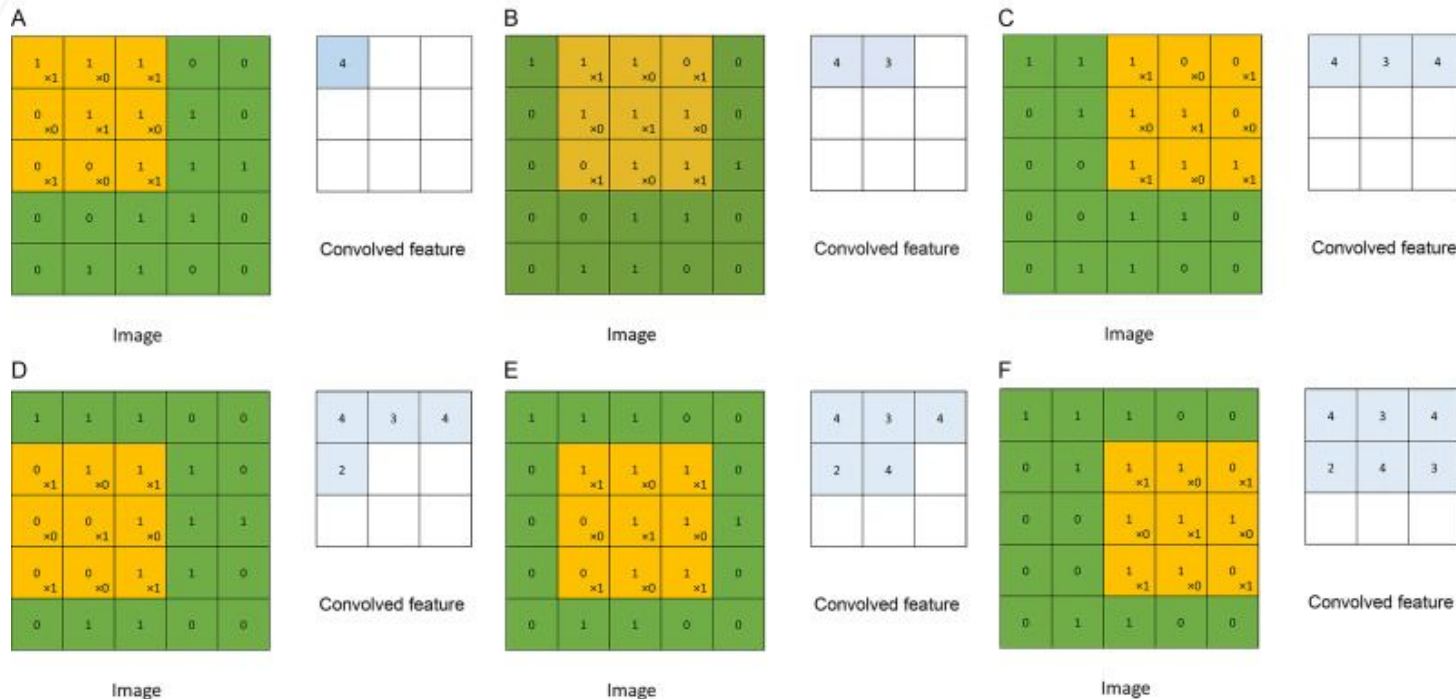
# Coding

# Part - 1: CNN from Scratch

- Modules
  - Convolution
  - MaxPooling
  - ReLU
  - Linear
  - Tie everything together: Convolutional Classifier
- Implement Optimizer:
  - SGD with momentum


# Part - 1: CNN from Scratch

## Convolution:



# Part - 1: CNN from Scratch

## Convolution:



# Part - 1: CNN from Scratch

## Convolution:

- Calculate final size of output
  - $H = 1 + ((\text{input.H} - \text{kernel\_size} + 2 * \text{padding}) / \text{stride})$
  - $W = 1 + ((\text{input.W} - \text{kernel\_size} + 2 * \text{padding}) / \text{stride})$
- Can use loops to iterate over batch, output\_channels, row, and column
- Use np.pad, np.dot



# Part - 1: CNN from Scratch

## Vectorization Tricks:

1. “Place” the kernel on the image
2. Multiply the kernel weights by the corresponding pixels
3. Sum everything together
4. “Stride” (or move) the kernel x amount of pixels
5. Repeat for the whole image
6. Repeat for all the kernels

# Part - 1: CNN from Scratch

## Vectorization Tricks:


$$\underbrace{\begin{bmatrix} 0 & 1 & 2 \\ 3 & 4 & 5 \\ 6 & 7 & 8 \end{bmatrix}}_{\text{Initial input}} \xrightarrow{\text{Divide in windows}} \left[ \begin{bmatrix} 0 & 1 \\ 3 & 4 \\ 6 & 7 \end{bmatrix} \quad \begin{bmatrix} 1 & 2 \\ 4 & 5 \\ 7 & 8 \end{bmatrix} \right] \cdot \underbrace{\begin{bmatrix} 0 & 1 \\ 2 & 3 \end{bmatrix}}_{\text{Kernel}}$$

Multiply with kernel

$$\left[ \begin{bmatrix} 0 & 1 \\ 6 & 12 \end{bmatrix} \quad \begin{bmatrix} 0 & 2 \\ 8 & 15 \end{bmatrix} \right] \xrightarrow{\text{Sum window values}} \underbrace{\begin{bmatrix} 19 & 25 \\ 37 & 43 \end{bmatrix}}_{\text{Final output}}$$


# Part - 1: CNN from Scratch

Vectorization Tricks: `np.lib.stride_tricks.as_strided`



# Part - 1: CNN from Scratch


## Vectorization Tricks:



# Part - 1: CNN from Scratch

## Vectorization Tricks:

(original\_row\_stride \* layer\_stride,  
original\_column\_stride \* layer\_stride,  
original\_row\_stride,  
original\_column\_stride)




# Part - 1: CNN from Scratch

## Vectorization Tricks:

```
new_shape = (out_height, out_width, kernel_size, kernel_size)  
new_strides = (rows_stride * layer_stride, columns_strides * layer_stride,  
rows_stride, columns_strides)
```

```
windowed_input = np.lib.stride_tricks.as_strided(input, new_shape, new_strides)
```




# Part - 1: CNN from Scratch

## Vectorization Tricks:

```
new_shape = (out_height, out_width, kernel_size, kernel_size)
new_strides = (rows_stride * layer_stride, columns_strides * layer_stride,
rows_stride, columns_strides)
```

```
windowed_input = np.lib.stride_tricks.as_strided(input, new_shape, new_strides)
```



# Part - 1: CNN from Scratch

Vectorization Tricks:

**np.einsum:**

```
windowed_input = np.random.rand(batch_size, channels, height, width, kernel_size,  
kernel_size)
```

```
weights = np.random.rand(num_filters, channels, kernel_size, kernel_size)
```

```
np.einsum( ? , windowed_input, weights)
```



# Broadcasting

$A = (a,b,c)$   $B=(b,)$

$A+B[:,\text{np.newaxis}]$

$A+B[:,\text{None}]$

Broadcast across 1st and 3rd dimensions before adding

# Unravel Index

```
arr_2d = np.array([[1, 2, 3],  
                  [4, 5, 6],  
                  [7, 8, 9]])
```

```
flat_index = 5
```

```
multi_index = np.unravel_index(flat_index, arr_2d.shape)
```

```
(1,2)
```

# Part - 1: CNN from Scratch

- Max-Pool
- Use loops to calculate the maximum value
- `np.unravel_index`


## Part-2

### Complete the training loop of PyTorch

- Implement the training and validation step utilizing PyTorch
- Key concept is that in validation you won't be updating the gradients


## Part-2

- Implement a basic CNN utilizing PyTorch



## Part-2

- Two fully connected layers with a sigmoid activation function in between



## Part-2: Your Model

- Your accuracy on test data needs to be above 0.5 to receive partial credits. For full credit you need above 0.8.
- You must upload the checkpoint because gradescope won't be able to train your model
  - The script will create and save the checkpoints for you.
- If you make your model too deep then gradescope might timeout as there is a limit for runtime + size (10 min for all tests, 100 mb filesize)
- You can utilize your own model selection. Take inspiration from architecture design discussed in lecture.